

STAMPED principles for reproducible research objects

Austin Macdonald¹ , Cody C. Baker¹ , Isaac To¹ , Yaroslav O. Halchenko¹ 

¹Center for Open Neuroscience, Department of Psychological and Brain Sciences, Dartmouth College, NH, USA 03755

August 2026

1 Abstract

Scientific claims increasingly rely upon the interplay of code, data, and computational environments. Yet the record of how they are used together is often incomplete, scattered, or lost. This undermines rigor, reproducibility, reusability, and efficiency. Previous approaches have improved the governance of digital objects but do not specify how research objects ought to be structured and managed so they can be re-executed and extended. The community is missing a shared vocabulary for this operational layer. Building on recurrent practices and convergent patterns across disciplines, we formalize seven principles: Self-containment, Tracking, Actionability, Modularity, Portability, Ephemerality, and Distributability (STAMPED). Together they give researchers guidelines for building and assessing research objects that others can trust, rerun, and build upon. We frame each principle as a spectrum so that adoption is incremental and starts from existing practice. We support each principle with normative requirements, an interactive checklist, and a map of enabling tools. As conventions mature, tooling improves, and AI agents become increasingly involved in research workflows, the goals of rigor, reproducibility, reusability, and efficiency are becoming more attainable. STAMPED gives researchers, collaborators, reviewers, and agents a common language, concrete goals, and an aligned direction for making computational research more durable.

2 Introduction

Data-driven research depends on specific versions of datasets, software, and computational environments, together with a record of how they were used to produce results. We adopt the term **research object** to denote a collection of data, code, and metadata that together represent the research as a complete unit. It is common for the components of a research object to be organized independently with code in one repository, data on a shared drive, environment setup in a wiki page, and provenance in a lab notebook. This separation undermines rigor, reproducibility, reusability, and efficiency. Even when its components are gathered in one place, a research object is often tightly coupled to specific software versions and environments which may be opaque to outside collaborators, making it difficult to reuse or redistribute.

The FAIR principles¹ and FAIR4RS guidelines² provide structured guidance for making digital objects "Findable, Accessible, Interoperable, and Reusable". The Workflow Community Initiative (WCI-FW)³ mapped these principles onto computational workflows but explicitly chose not to define new ones, instead treating workflows as hybrid data & software objects. STAMPED addresses a complementary layer: the day-to-day practices that underpin re-execution, extension, and redistribution of a research object. Both the FAIR and STAMPED frameworks cover overlapping concerns. For example, persistent identifiers serve both Findability and Tracking but accent different aspects. FAIR emphasizes discovery and interoperability across repositories, while STAMPED emphasizes improving the operational maturity of scientific workflows in day-to-day research practice.

Many researchers already employ practices that address these operational properties, such as version control, containerized environments, structured project directories, and documented analysis steps. Yet the community lacks a shared vocabulary for referring to these ideas, evaluating how thoroughly they are applied, and communicating that evaluation to collaborators and reviewers. Here we provide that

vocabulary by introducing seven principles which articulate how Self-contained, Tracked, Actionable, Modular, Portable, Ephemeral, and Distributable (STAMPED) a research object is, thus describing its operational maturity. These originate from both the YODA Principles⁴, as well as VAMP⁵. Rather than a compliance threshold, each STAMPED principle is a spectrum from a practical minimum—often what researchers are already doing—to an aspirational ideal. A researcher who stores everything under one project directory, version-controls analysis scripts, and documents how to run them is already meeting the practical minimum for Self-containment, Tracking, and Actionability. As complexity grows, such as for multi-institutional collaborations with heterogeneous workflows spanning platforms enforcing long-term reproducibility, these same principles guide researchers toward more rigorous practices without requiring wholesale adoption of new tools. STAMPED provides a descriptive vocabulary and a framework for evaluation that can incrementally improve the operational properties of research objects towards enhanced rigor, reproducibility, reusability, transparency, and efficiency.




Figure 1: Illustration of the STAMPED principles.

3 The STAMPED Principles

Table 1: The STAMPED Principles where each one addresses a distinct property of research object organization that contributes to reproducibility. Together, they provide a framework for evaluating and improving the reproducibility of computational research.

Letter	Principle	Description
S	Self-containment	A research object is a complete retrieval unit; it can be obtained and understood in its intended scope without needing to reference external resources.
T	Tracking	The state and provenance of all components is recorded.
A	Actionability	Research object contains machine-actionable information to carry out procedures to obtain or reproduce content.
M	Modularity	All modules are independent and composable.
P	Portability	Research object could be moved to different environments while retaining its STAMPED properties.
E	Ephemerality	Procedural execution is performed within a throwaway environment.
D	Distributability	All modules and procedures are shareable externally in a persistent state.

The formal requirements use the graded keywords of RFC 2119⁶: MUST defines the practical minimum, while SHOULD and MAY indicate progressively more aspirational practices.

3.1 Self-containment

- **S.1:** All modules and components essential to replicate computational execution **MUST** be reachable within a single top-level research object.
- **S.2:** License declarations for all components are themselves part of the research object and **MUST** be retrievable alongside the content they govern.

Reproducibility fails when a research object depends on resources that are not part of it; an undocumented library, a dataset referenced by a broken URL, a script that assumes tools are installed on the host system. We call this the “don’t look up” rule: a research object must never rely on any implicit external state. Instead, it must be a complete retrieval unit, where all modules and components essential to reproduce its results are contained within a single top-level boundary (S.1, Table 1).

Components may be included literally (e.g., files committed directly) or by reference (e.g., as subdatasets, registered data URLs). Both approaches satisfy self-containment, provided the references are explicit and part of the research object. At a minimum, everything needed is gathered under one root and any external dependency is explicitly documented. This ensures that nothing is implicitly assumed about the host system, a concern further addressed by Portability (P). At the ideal end, every reference is precise enough that retrieval is unambiguous; how to achieve that precision is the concern of Tracking (T) and Distributability (D).

Self-containment is the foundation upon which the remaining STAMPED principles are built. Tracking (T), Modularity (M), Portability (P), and Distributability (D) each address a different aspect of how that boundary is organized, versioned, and shared—but none are meaningful without first establishing what is inside it.

3.2 Tracking

- **T.1:** Persistent content identification **MUST** be recorded for all components.
- **T.2:** All components **SHOULD** be tracked using the same content-addressed version control system.
- **T.3:** Provenance of all modifications **MUST** be recorded.
- **T.4:** Code-driven provenance **SHOULD** be recorded programmatically and **MUST** include the versions of all components involved.

Research data and code change constantly during a project, and each change compounds the difficulty of reconstructing any previous state, such as which version of a dataset was used for a particular analysis, what parameters were set, and what code produced a given figure. A Self-contained (S) research object gathers everything needed in one place, but without a record of how each component arrived at its current state, reproducing a prior result becomes guesswork. Tracking addresses two concerns: 1) the identity of each component and its exact state at a point in time, and 2) **provenance**, the recorded history of the actions by which identities produced or modified it. These are often managed separately, which makes validating past work harder.

The spectrum of tracking begins with content-addressed version control. Rather than relying on semantic version labels, content-addressed systems identify each state by a checksum of its contents. This distinction matters; two datasets labeled “version 1.0” by different labs are ambiguous whereas two datasets with identical content hashes are provably identical. Version control commits also capture basic provenance of who made the change, when, and a description of why. Cryptographic signatures (e.g., signed git commits and tags, sigstore-signed releases) strengthen these records from self-reported to verifiable, allowing recipients to confirm that components and changes originate from the claimed identities. This unifies identity and provenance in a single system and also provides a safety net since any change can be reverted and any prior state restored, making it easier to modify and experiment. For tracking to be effective, all components of the research object must be under version control: code, data, environment definitions, and configuration.

Traditional version control is designed for code and other text files, but research objects often include large binary datasets, container images, and other assets that do not fit this model. Compatible tools such as git-annex, DataLad⁷, or DVC⁸ extend content-addressed tracking to these components, ensuring that all parts of the research object are managed under the same system.

Toward the ideal end, provenance is captured programmatically by tooling rather than by manual annotation. Each computational step should record exactly what command was executed and what inputs were consumed, producing richer and more reliable records than a human would typically write. When modules are independently tracked under compatible systems, the version of every module at the time of each commit is also discoverable, providing a complete picture of the state that produced a given result. This enables not just inspection but re-execution. Beyond the research object itself, knowing the state of data and code at each hand-off between experimental, analytical, and engineering stages is what allows a team to coordinate updates without losing the provenance chain—an operational view of tracking developed by the SciOps framework⁹. Across communities, multiple ontologies and packaging standards have emerged to make these provenance records portable across tools, including cross-domain models such as W3C PROV¹⁰, packaging formats such as RO-Crate¹¹, and domain-specific specifications such as BIDS-Prov (BEP028)¹². The relationships among such provenance, assertion, and evidence ontologies have been surveyed recently¹³, emphasizing the already available rich ecosystem to be promoted and wider adopted for improving tracking of scientific processes.

3.3 Actionability

- **A.1:** Research object **MUST** contain sufficient instructions to reproduce all computational results.
- **A.2:** Procedures **SHOULD** be specified as executable specifications.

A research object may be Self-contained (S) and fully Tracked (T), yet having all ingredients and a record of what was done is insufficient without a recipe that can be followed. Requirement A.1 is what makes a research object operationally useful by bridging the gap between having the components and being able to use them.

At a minimum, a research object must contain documentation thorough enough for another researcher to follow every step to reproduce results. Actionability applies to every other STAMPED principle. At the ideal end, each principle is enacted not through documentation alone but through executable specifications (A.2):

- **Self-containment** is more actionable when retrieval of all components is specified and executable (e.g., `git clone`, `datalad install`), not just listed in a manifest.
- **Tracking** is more actionable when provenance records are re-executable (e.g., `datalad rerun`), not just inspectable.
- **Modularity** is more actionable when components can be composed, updated, and replaced via tooling (e.g., `git submodule`), not just organized into directories.
- **Portability** is more actionable when environment specifications are machine-readable and can be resolved on different hosts (e.g., `conda env create`, Nix flakes), not just documented in a README.
- **Ephemerality** is more actionable when computation can be orchestrated in disposable environments (e.g., `docker compose`, Slurm), not just instructed to “run in a clean environment.”
- **Distributability** is more actionable when a frozen state can be produced and retrieved by others (e.g., containers, `npm ci`), not just described with pinning instructions.

The principle is tool-agnostic; any system that moves a property from documented to executable moves the research object further along the actionability spectrum. Actionability also extends from tooling to the operational layer of a research group. A SciOps-style workflow makes hand-offs between experimental, analytical, and engineering stages explicit so that artifacts move from one step to the next with unambiguous state transition⁹, an organizational counterpart to the per-component actionability described above.

3.4 Modularity

- **M.1:** Components **SHOULD** be organized in a modular structure.
- **M.2:** Components **MAY** be included directly or linked as subdatasets.

- **M.3:** Each module’s license SHOULD be declared independently and checked for compatibility at the boundaries where modules are combined

A research object may be Self-contained (S) and Tracked (T), yet when all components are interleaved in a single monolithic directory, it can be fragile and difficult to navigate. Separation of concerns is a foundational engineering practice: organizing a system so that each piece has a clear, distinct role. With clear boundaries between the conceptual components of a research object, it becomes more navigable and maintainable. We define a **module** as a separately Distributable (D) unit of a research object. This can include a subdataset, a collection of analysis scripts, or the definition for some computational environment. In practice, module boundaries often align with repository boundaries, but a module may also be a published container, a registered dataset, or any independently versioned unit. Modularity applies this principle to research objects, structuring them as assemblies of modules whose boundaries are explicit. Module boundaries are also where license compatibility must be checked, since combining modules with incompatible terms can prevent the assembled research object from being lawfully redistributed.

At a minimum, a research object should separate analysis code, input datasets, computational environments, licenses, configuration settings, and results into distinct directories. Beyond this, independent versioning of modules enables reuse; a dataset or environment can be updated to a newer version, or swapped for an alternative, without disrupting the rest of the research object. This is possible through manual management but becomes practical when handled by tooling that automates module installation, versioning, and composition. Toward the ideal end, the full research object can be reassembled from its specification through automated, recursive installation of its modules.

3.5 Portability

- **P.1:** Procedures MUST NOT depend on any undocumented host environment state.
- **P.2:** Computational environments MUST be explicitly specified.
- **P.3:** Environment definitions MUST be version controlled.

A research object that is Self-contained (S), Tracked (T), and Modular (M) may still fail to run on a different machine if it is silently affected by any inherited host state. Common examples of such states include specific developer-level package installations with no released versions, unique system libraries exposed to custom user profiles, hardcoded system paths, or an assumed OS configuration (e.g., path character limits). This coupling is often invisible on the original system and only surfaces when someone else attempts to reproduce the work. Portability requires making all host dependencies explicit, so that the research object can be adapted to a new environment.

Some degree of coupling to the host is unavoidable, such as path naming restrictions (“C:/” vs. “/mnt”), resource limits (minimum CPU/RAM requirements), or platform-specific flags. The spectrum of portability begins with isolating this host-specific configuration from the rest of the research object, so that adapting to a new environment, however involved that may be, requires only changing the configuration. Beyond this, several complementary approaches reduce host coupling. Virtual machines provide full system-level isolation; container-based tools (Docker, Singularity/Apptainer) bundle OS-level environments; declarative package managers (Nix¹⁴, Guix¹⁵) specify environments that can be reconstructed from a definition. Toward the ideal end, the computational environment is fully specified and version-controlled within the research object, enabling it to be instantiated on any compatible system.

3.6 Ephemerality

- **E.1:** Computational results SHOULD be produced in ephemeral environments.

A research object may satisfy Self-containment (S), Tracking (T), and Portability (P) while its computational environment has only ever been assembled once, incrementally, on the original researcher’s machine. During development, researchers routinely modify their environment by installing packages, adjusting configurations, and setting variables. Each change is individually small but collectively difficult to document manually. The result is environment drift, where an environment that works but may not be reconstructable from its

specification alone. When this occurs, undocumented steps can go unnoticed when internalized knowledge silently fills the gaps. Ephemerality eliminates this problem by construction; a disposable environment built from Tracked (T) specifications cannot carry any undocumented state, so that successful execution exercises Self-containment (S), Actionability (A), and specification completeness.

The spectrum of ephemerality ranges from manually setting up a fresh environment from the project’s specification (which is usually sufficient to catch environment drift) to fully automated disposable environments created and destroyed per execution. At the ideal end, ephemeral computation also enables testing across different host systems, where same-host execution validates specification completeness and different-host execution validates that specifications are not coupled to a specific system. The FAIRly big workflow¹⁶ and other approaches to using Distributed (D) computing platforms such as SLURM or Code Ocean¹⁷ operationalize this ideal, treating each computation as an ephemeral work unit. This approach easily enables horizontal scaling, as independent disposable instances can be parallelized across subjects, parameters, or datasets.

Yet ephemeral computation does not validate everything. For instance, an ephemeral environment with network access may fail to notice unpinned dependencies fetched at build time, a Self-containment (S) gap that only surfaces when those resources move or disappear. Nonetheless, the principle substantially raises confidence that the research object is Self-contained (S), Portable (P), and Actionable (A).

3.7 Distributability

- **D.1:** All referenced modules and components **MUST** be persistently retrievable by others.
- **D.2:** Environment specifications **SHOULD** support reproducible builds.
- **D.3:** Each module **SHOULD** carry an explicit license with a resolvable identifier.

A research object that is Self-contained (S) on the author’s machine but cannot be obtained by others in the same state is effectively unreproducible beyond its origin. Self-containment (S) establishes that everything needed is within the boundary and Tracking (T) records the state of each component. Distributability ensures that others can obtain the research object in that same state.

The distinction mirrors the concept of a software distribution: a curated, versioned bundle in which all components are resolved to specific versions and packaged for consumption. Simply sharing a research object by uploading scripts to a repository with loose dependencies, or posting files on a website with no version or other persistence guarantees, does not constitute distribution in this sense. A distributable research object is packaged so that it can be retrieved in the same state as intended.

Distributability also has a circular relationship with the other STAMPED principles in that someone else’s distribution effort often serves as the starting point for a new research object’s Self-containment (S). Researchers routinely begin by downloading containers, fetching published datasets, and installing released software—Modularly (M) composing the distributions of others to assemble their own Self-contained (S) research objects. Though persistent retrievability is a necessary but not sufficient condition for redistribution. Additionally, each module should carry an easily identifiable license or set of licenses for any bundled assets (e.g., via SPDX identifiers¹⁸ or the REUSE specification¹⁹) so that downstream researchers can compose, modify, and re-share components without legal ambiguity.

The spectrum of distributability begins where FAIR leaves off, with publicly accessible components and documented retrieval instructions. Toward the ideal end, several complementary strategies reduce the risk that external changes will break a research object’s reproducibility. Bundling dependencies into containers or archives reduces the surface area of components that can independently change or disappear; hosting on archival infrastructure (Zenodo, DANDI, Software Heritage) with content-addressed identifiers allows recipients to verify integrity regardless of source; mirroring across multiple platforms protects against the failure of any single registry. Where downstream consumers must verify not only integrity but also origin and build process, signed releases and supply-chain attestations (e.g., sigstore, in-toto, SLSA) provide cryptographic evidence that an artifact was produced by the claimed identity from the claimed sources. This becomes important when distributed artifacts will execute in environments enforcing code-signing policies, such as notarization or signed container admission.

Table 2: Normative statements about STAMPED principles as they apply to research objects.

Principle	Requirements
To be Self-contained:	• all modules and components needed to understand and execute the Research Object are retrievable as a single unit • external dependencies are explicitly documented with retrieval instructions • there are no implicit references to undocumented external resources
To be Tracked:	• every component has version information (commit hash, tag, or identifier) • changes to components are recorded with timestamps and authorship • provenance records capture the computational history, context, and transformations
To be Actionable:	• instructions for executing procedures are present and unambiguous • execution paths can be followed manually or automated programmatically • the Research Object transitions from documentation to operational capability
To be Modular:	• modules can be independently Tracked (T), modified, and Distributed (D) • components are organized in logical, separable units • modules can be composed together or used in isolation
To be Portable:	• system requirements and dependencies are explicitly documented • the Research Object is flexible enough to execute on different host environments without modification by the user • environment specifications are machine-readable where possible
To be Ephemeral:	• computation can occur in temporary, disposable environments • results are reproducible without knowledge of previous runs • the environment is rebuilt from its specification before each execution, ensuring no reliance on undocumented state
To be Distributable:	• all modules and components can be shared in a persistent, retrievable state • dependencies are frozen or pinned to specific versions across systems • the Research Object can be obtained by others in the same state as intended

3.8 Checklist for compliance to principles

We hope that the preceding sections have provided a clear understanding of the STAMPED principles and their rationale. While this understanding is essential, researchers may also benefit from a practical checklist to assess their research objects against these principles. Similar to the living examples repository, an interactive checklist has been provided at the checklist.stamped-principles.org webpage to guide assessments of compliance with STAMPED principles, ordered by the strength of the requirement (MUST, SHOULD, MAY). The schema for this checklist is itself a living, version-tracked entity²⁰ that is likely to grow over time as more common violations of the principles are discovered in practice.

3.9 Enabling tools

No single tool addresses all the dimensions of scientific rigor described above, though many tools contribute to one or several properties simultaneously. For example, container technologies such as Docker or Singularity serve Portability (P) by providing OS-level isolation from explicit environment specifications, enable Ephemerality (E) by disposing per-execution environments, and facilitate Distributability (D) by producing frozen, shareable artifacts. Similarly, version control systems such as Git underpin both Tracking (content-addressed identification and change history) and Self-containment (gathering all components under a single root).

In practice, Self-containing (S) all data by storing it directly within the research object may be undesired or prohibitive due to size, licensing, or privacy concerns. This creates a tension with Tracking (T), which requires that every component be identified and versioned. Extensions such as git-annex, DataLad, DVC, and Git LFS strike the balance between the two principles by tracking metadata that identifies specific data versions and obtainability information instead, e.g., content checksums and URLs of remote servers that act as a source, and providing Actionable (A) interfaces to obtain the data when necessary.

This overlap is not incidental — it reflects the interdependence of the STAMPED principles themselves. Achieving Ephemerality (E), for instance, inherently exercises Portability (P) and Actionability (A); a disposable environment must be reconstituted from explicit specifications, encouraging end-to-end automation. Likewise, Distributability (D) builds on Self-containment (S) and Tracking (T); a research object can only be shared in a consistent, retrievable state if its boundary is well-defined and all components are precisely identified. The tools that enhance multiple properties of a research object do so precisely because they operationalize these connections. Table 3 maps some commonly used tools and technologies to each STAMPED principle.

STAMPED principles are tool-agnostic. A well-chosen tool can improve several properties at once, and a thoughtfully assembled toolchain can cover all of them. Researchers should select tools based on their domain, infrastructure, and team familiarity, recognizing that the tools listed here are illustrative rather than prescriptive. What matters is that the underlying property is satisfied, regardless of which tool achieves it. Where multiple tools enhance the same STAMPED property of an object, the choice between them often involves trade-offs between flexibility and simplicity.

A representative example is the choice between git-annex and Git LFS for large-file tracking. Both appear under Self-containment (S) and Tracking (T) in Table 3, but they occupy different points on a complexity-flexibility spectrum. git-annex supports many remote types (e.g., other git repositories over SSH or HTTP, or simple cloud data stores such as S3, local drives, offline USB archives), making it well-suited to multi-institutional collaboration, long-term archival where data sovereignty or offline access is required, as well as local content tracking. Git LFS offers a simpler alternative; it is transparent to users, natively supported by GitHub and GitLab, and broadly adopted. But it is centralized (requiring an LFS server), lacks offline capability, and provides less flexible remote configurations. Lightweight STAMPED implementations using Git LFS are entirely feasible for easier onboarding, trading flexibility for reduced operational overhead. The same pattern of flexibility versus simplicity recurs across other tool choices in Table 3: container-based versus package-based environment management, workflow engines versus shell scripts, and distributed archives versus centralized platforms.

MUST S1 Self-containment 1/3

All modules and components essential to replicate computational execution MUST be reachable within a single top-level research object.

Are all files and directories nested under a common root? ☒ Yes ☐ No

Are datasets included, linked, or referenced reachable from the code and environment specifications without crossing the boundary of a common root? ☐ Yes ☒ No

Reason for not meeting requirement...

Is external software included in the environment or linked as submodules within the project boundary? ☐ Yes ☐ No

MUST T1 + T3 + T5 Tracking 0/3

Persistent content identification MUST be recorded for all components. Provenance of all modifications MUST be recorded. Provenance records MUST include the versions of all components involved.

Are version control systems such as `git` used for code, test, documentation, and configuration files? ☐ Yes ☐ No

Are version control systems such as `git-annex`, `DataLad`, or `git LFS` used for large binary data? ☐ Yes ☐ No

Are the exact environment specifications used to generate a set of results included in the provenance records to link computational actions to particular environments? ☐ Yes ☐ No

MUST A1 Actionability 0/2

Research object MUST contain sufficient instructions to reproduce all computational results.

Is a `README.md` or `Makefile` included with instructions for installation and usage? ☐ Yes ☐ No

Is there a clear starting point for users to start reproducing results (e.g., a main script, a workflow definition, or a container image)? ☐ Yes ☐ No

MUST P1 Portability 1/3

Procedures MUST NOT depend on undocumented host environment state.

Are relative paths used in scripts, avoiding hardcoded or system-specific paths like `C:\Users\...` or `/home/user/...`? ☐ Yes ☐ No

Are all software dependencies included in the environment specification, rather than relying on pre-installed tools? ☒ Yes ☐ No

Have all assumptions about the host system's configuration been documented (e.g., specific OS versions, required system libraries, or environment variables)? ☐ Yes ☐ No

MUST P2 Portability 0/1

Computational environments MUST be explicitly specified.

Is there a clear list of system requirements and dependencies documented in the `README` or environment specifications? ☐ Yes ☐ No

MUST P3 Portability 0/2

Environment definitions MUST be version controlled.

Are environment specifications (e.g., `Dockerfiles`, `pyproject.toml`, `package.json`) included in the version control system alongside code and data? ☐ Yes ☐ No

Is there a process for updating environment specifications when dependencies change, and are these updates tracked in version control? ☐ Yes ☐ No

MUST D1 Distributability 1/3

All referenced modules and components MUST be persistently retrievable by others.

Are environment specifications (e.g., container digests, frozen package manifests) included to ensure others can attempt to exactly replicate the environment? ☒ Yes ☐ No

Are environment specifications shared in a way that others can access and use them (e.g., published container images, archived environment files)? ☐ Yes ☐ No

Is there documentation on how to obtain and use the environment specifications for reproduction? ☐ Yes ☐ No

SHOULD T2 Tracking 0/1

All components SHOULD be tracked using the same content-addressed version control system.

Is a common version control system (e.g., `git`, `DVC`, `git-annex`, or `DataLad`) used across all components? ☐ Yes ☐ No

SHOULD A2 Actionability 0/1

Procedures SHOULD be specified as executable specifications.

Is the workflow tested regularly to ensure instructions remain accurate? ☐ Yes ☐ No

SHOULD M1 Modularity 0/1

Components SHOULD be organized in a modular structure.

Are raw data, processed data, code, and environment definitions separated into distinct modules? ☐ Yes ☐ No

SHOULD E1 Ephemerality 0/2

Computational results SHOULD be produced in ephemeral environments.

Is the pipeline tested in a fresh container, batch job, clean virtual machine, or cloud-based instance? ☐ Yes ☐ No

Does the workflow spawn disposable environments for each execution, allowing for isolation of runs? ☐ Yes ☐ No

SHOULD D2 Distributability 0/3

Environment specifications SHOULD support reproducible builds.

Are the exact environment specifications used to generate a specific set of results regularly tested to ensure they can be reconstituted as intended? ☐ Yes ☐ No

Are environment artifacts archived within the research object where possible (e.g., executable binaries, container images)? ☐ Yes ☐ No

Is there a process for updating environment specifications when dependencies change, and are these updates tracked in version control? ☐ Yes ☐ No

MAY M2 Modularity 0/5

Components MAY be included directly or linked as subdatasets.

Are external datasets or software dependencies included as submodules or linked as submodules? ☐ Yes ☐ No

Is the modular structure documented to clarify how components relate and can be recombined? ☐ Yes ☐ No

Are modular boundaries defined in a way that allows independent updates without breaking the overall workflow? ☐ Yes ☐ No

Is there a clear mechanism for composing modules together (e.g., `git` submodules, or container orchestration)? ☐ Yes ☐ No

Is there documentation or tooling to support users in understanding and navigating the modular structure? ☐ Yes ☐ No

Figure 2: Screenshot of the interactive checklist for helping guide a workflow developer through common cases.

Table 3: Sample tools to improve scientific practice for each STAMPED principle.

Principle	Tools / Technologies	Role
Self-containment (S)	Git submodules, DataLad, git-annex, DVC, Git LFS	Gather all components (code, data, environments) under a single root via direct inclusion or explicit references (e.g., subdatasets, registered URLs)
Tracking (T)	Git, git-annex, DataLad, DVC, Docker and other containers, Kedro, Git LFS, OSF.io, Zenodo and other data archives providing DOIs	Version control for code and data; content-addressed identification; provenance recording (e.g., <code>datalad run</code> for programmatic provenance capture); persistent identifiers for content (DOIs)
Actionability (A)	Make, Snakemake, Nextflow, CWL, <code>datalad run</code> , <code>git annex addcomputed</code>	Define executable specifications for workflows; provide clear entry points for reproducing results
Modularity (M)	Git submodules, DataLad subdatasets, Kedro	Organize components into independently versioned, composable modules; provide project templates with logical directory separation
Portability (P)	Docker, Singularity, Nix, conda, pip (<code>pyproject.toml</code> , <code>requirements.txt</code>), npm	Explicitly specify and isolate computational environments; container-based (OS-level isolation) or package-based (declarative, reproducible builds)
Ephemerality (E)	Docker, Singularity, Docker Compose, Slurm, cloud deployments (e.g., AWS, GCP, Azure, GitHub Actions)	Spawn disposable, temporary environments per execution; validate that specifications are complete by reconstituting from scratch
Distributability (D)	Zenodo, PyPI, conda-forge, GitHub Release CDN, DockerHub/GHCR, RO-Crate, <code>npm ci</code> , <code>pip freeze</code> , executable compilers	Archive and persistently host frozen, versioned research objects and their components; enable retrieval by others in the intended state

3.10 Examples and Case Studies

3.10.1 Pedagogical examples

A comprehensive up-to-date review of existing tools is out of scope for this formalization paper and would quickly become outdated. To fill that niche, we initiated a satellite project `stamped-examples`²¹ with a website to describe tools and examples, and characterize them in terms of STAMPED and FAIR principles at the examples.stamped-principles.org website. These complement the formal definitions and walk a reader through the STAMPED properties using concrete examples. The pages tagged “STAMPED-intro” (examples.stamped-principles.org/tags/stamped-intro/) trace how each property applies to the same simple research object, providing a gentle on-ramp before the real-world adoptions below.

3.10.2 Real-world adoptions

Several research projects have already adopted and documented the YODA principles (the predecessor to STAMPED) demonstrating practical utility across domains.

BIDS Standard Evolution: A significant validation of the “do not look up” principle occurred within the Brain Imaging Data Structure (BIDS)²² community. Originally, the BIDS specification nested derived data under `derivatives/` within the original raw dataset, creating upward dependencies that violated Self-containment (S) and Modularity (M), which broke Portability (P). Following YODA principles, the BIDS community reversed this relationship so that derivative datasets now exist independently and reference raw data as subdatasets, not vice versa.

This architectural shift influenced major neuroimaging tools:

- **fMRIPrep**²³: Switched default output layout to follow YODA structure
- **OpenNeuroDerivatives**²⁴: Entire derivative dataset collection now follows YODA organization, separating processed outputs from raw data dependencies
- **BIDS specification**: Updated to accommodate and recommend YODA-compliant layouts for derivative datasets

Workflow Platforms: Infrastructure projects implementing YODA at scale:

- **BABS**²⁵: Platform implementing the “FAIRly big workflow” approach, demonstrating YODA principles for large-scale neuroimaging analyses with containerized pipelines and modular dataset composition
- **CRCNS.org**: Neuroscience data sharing platform providing YODA-structured templates and validation tools

These adoptions demonstrate that STAMPED principles solve practical organizational challenges across scales, from individual studies to community-wide infrastructure.

3.11 Convergent evolution of data-driven scientific practices

The challenges that motivate STAMPED are not unique to any single domain.

Fragmented provenance, irreproducible computations, and ad-hoc project layouts have plagued communities across geophysics, genomics, neuroimaging, statistics, and machine learning for three decades. Independently, each converged on strikingly similar organizational responses. What follows is not a comprehensive review but a tracing of that convergence, organized chronologically, to show that the principles STAMPED formalizes reflect durable responses to recurring challenges.

The earliest calls for reproducibility-as-organization appeared in the 1990s. Claerbout and Karrenbach²⁶ demonstrated that coupling electronic documents with Makefiles could make geophysical analyses self-contained, tracked, and re-executable—anticipating Self-containment (S), Tracking (T), and Actionability (A). Buckheit and Donoho²⁷ extended this idea with WaveLab, a compendium packaging code, data, and figures so that “an article ... is not the scholarship itself, ... the actual scholarship is the complete software development environment and the complete set of instructions which generated the figures”—touching Self-containment (S), Actionability (A), Portability (P), and Distributability (D). Gentleman and Temple Lang²⁸

formalized the research compendium concept, emphasizing modular packaging for reuse and redistribution—Self-containment (S), Modularity (M), Actionability (A), and Distributability (D). Noble²⁹ provided the first widely adopted project-organization guide for computational biology, offering concrete directory-structure conventions that addressed Self-containment (S) and Modularity (M).

A wave of normative frameworks followed in the 2010s. Stodden³⁰ articulated legal and social dimensions of sharing computational research, raising concerns central to Distributability (D) and Tracking (T). Sandve et al.³¹ distilled ten rules for reproducible computation (track versions, record workflows, archive environments)—spanning Tracking (T), Actionability (A), and Self-containment (S). The FAIR principles¹ established findability, accessibility, interoperability, and reusability as community norms for digital objects, later extended to research software by FAIR4RS² and to workflows by the Workflow Community Initiative³. Wilson et al.³² codified “Good Enough Practices” that operationalized many of these ideals for working scientists—Actionability (A), Tracking (T), and Self-containment (S). The YODA principles⁴ combined these insights into organizational practices for research data management. De Smedt et al.³³ articulated the notion of FAIR Digital Objects as “actionable knowledge units,” foregrounding Actionability (A) as a first-class concern complementary to FAIR. Building on YODA and this emphasis on actionability, the VAMP formulation⁵ (Versioned, Actionable, Modular, Portable) made Actionability (A), Portability (P), and Modularity (M) explicit organizational requirements, and introduced Ephemerality (E) as validation for Portability (P)—directly prefiguring STAMPED. Notably, the WCI-FW framework³ extended FAIR to computational workflows but explicitly declined to define new principles, leaving the operational gap that STAMPED fills (how a research object should be structured and managed).

Independent of these normative efforts, a parallel “Git for data” convergence occurred in tooling. git-annex (2010), Git LFS (2015), DVC (2017), Pachyderm (2014), and Icechunk³⁴ (2024) all independently arrived at content-addressed identification of large files, logical separation of metadata from bulk storage, and derivation records linking inputs to outputs. These systems vary widely in architecture, ranging from cloud-native to local-first and from centralized to distributed. Despite these differences, each converged on the same core insight, that version-control semantics proven for code could and should extend to data. This convergence is evidence not of influence but of shared need, and maps most directly onto Tracking (T).

A complementary convergence emerged in platforms and metadata standards. Though not an exhaustive list: Galaxy³⁵ (first released 2005), Code Ocean¹⁷ (2017), and brainlife.io³⁶ each provided turnkey reproducibility services, each implementing separation of code, data, and environment as a first-class abstraction—reflecting Self-containment (S), Actionability (A), and Portability (P). On the standards side, the W3C PROV data model (2013), Frictionless Data³⁷ (2007), RO-Crate¹¹, and BioComputeObject³⁸ (IEEE 2791-2020, biocomputeobject.org) each addressed interoperable packaging and provenance recording. The facts that FAIR required elaboration to accommodate software² and workflows^{3;39}, and that independent standards converged on provenance and packaging, confirms that these dimensions are recurrently demanded by research practice.

Preregistration is the practice of committing to a study plan before outcomes are observed, which extends Tracking (T) to the intent behind a study. It was established in clinical trials by the early 2000s and gradually expanded into behavioral, social, and computational research⁴⁰. The Open Science Framework⁴¹ (2013) combined that practice with a “hub” for collaboration and storage, producing read-only snapshots of a study’s plan, data, code, and manuscripts under a single DOI, reflecting Self-containment (S) and Distributability (D). In some cases, pre-registration can also be considered a precursor to the Ephemerality (E) of the actual scientific pipeline where work is concentrated on formulating hypothesis, planing and potentially simulations, thus making actual later implementation of it replaceable.

Framework tooling and domain standards reveal the same pattern at the project-layout level. Cookiecutter Data Science (2016), BIDS²² (2016), Kedro (2019), nipoppy⁴² (2023), and others, all imposed opinionated directory structures and modular separation (independently motivated by the costs of ad-hoc layouts)—reflecting Modularity (M) and Self-containment (S). Each framework constrains where code and documentation reside, how raw data inputs and derived outputs are handled, which makes such projects navigable by *convention*.

Viewed together, these independent lines of development map onto every STAMPED principle. Self-containment (S) appears wherever projects bundle their dependencies into a single rooted structure; Tracking (T) wherever content-addressed storage or version control records the history of changes or acquisition of every component; Actionability (A) wherever Makefiles, workflow managers, or platform APIs make

re-execution or acquisition a single command; Modularity (M) wherever directory conventions or subproject composition enforce a separation of concerns; Portability (P) wherever good coding practice, containers or environment specifications decouple a research object from a particular machine; Ephemerality (E) wherever composition brings together computation on data which is treated as disposable and regenerable; and Distributability (D) wherever licensing, packaging standards, or data logistic mechanisms enable sharing.

The formalization presented here provides a shared vocabulary for what these practices have in common, enabling researchers to evaluate and incrementally improve their workflows regardless of the specific tools they employ. Figure 3 summarizes these parallel timelines and their mapping to STAMPED principles.

3.12 Future Directions

STAMPED principles provide a vocabulary rather than a closed specification, which is rooted in practices that continue to evolve with tooling, policy, and scientific convention. The directions below identify open fronts where the framework’s vocabulary will need to stretch, and where complementary work by adjacent communities will shape how STAMPED assists researchers in the coming decade. We begin with directions that apply broadly across computational science, including generalizing Ephemerality (E) to workflow specifications, aligning provenance standards, building adoption incentives, and seeding training pipelines. The remainder of the section then addresses the more specific and fast-moving challenges introduced by AI, where several STAMPED principles need explicit extensions and where the stakes for transparent, reproducible research are most immediate.

3.12.1 Specification-centric research objects as the reproducibility ideal

Ephemerality (E), as formulated here, treats the compute environment as disposable while the code and data it runs on remain durable. A complementary view is emerging in which workflow implementations themselves become regenerable from a durable specification. Declarative workflow languages such as CWL⁴³, Snakemake⁴⁴, and Nextflow⁴⁵, together with fully-declarative environment managers such as Nix¹⁴ and Guix¹⁵, already operate this way. The specification accumulates the domain knowledge, including inputs, constraints, expected outputs, and validation criteria, while any given implementation can be rebuilt or replaced without loss. Under this view, STAMPED principles apply to the specification first and the generated implementation second. The specification is what must be Self-contained (S), Tracked (T), Actionable (A), and Distributable (D), while the running code is Ephemeral (E) in the same sense as a container. The connection to pre-registration is direct and generalizes Ephemerality (E) even further⁴⁰. A well-formulated study specification is a commitment that survives any particular implementation, and what is Ephemeral (E) expands from the compute environment to the full “scientific middleware” (data acquisition, harmonization, preprocessing, analytics), with each stage reduced to a specification from which implementations can be derived. OSF Registrations⁴¹ already operate this way for preregistered plans, fixing a study’s specification under a DOI before any implementation begins. That is where concentration on formalization and standardization of overall study specifications becomes the target to apply STAMPED principles, so that the middleware steps could also be automatically carried out in an Ephemeral (E) fashion and potentially while exploring various possible implementations. The cost of under-specified middleware is already visible in practice; the NARPS study showed that 70 independent teams analyzing the same fMRI dataset with nominally similar hypotheses produced substantially divergent results⁴⁶, precisely because the study specification left most of the middleware up to each team.

3.12.2 Provenance standards convergence

As FAIR leaves specific protocols and standards open, STAMPED likewise does not prescribe a serialization for provenance records, and in practice researchers rely on whatever their tooling emits. Convergence among existing approaches would let STAMPED research objects carry Portable (P) provenance across toolchains without lock-in. Relevant efforts include cross-domain standards such as W3C PROV¹⁰ and RO-Crate¹¹, tool-level records such as DataLad’s⁷ `datalad run` provenance, and domain-specific specifications such as BIDS-Prov (BEP028)¹². This is less a change to STAMPED itself than a call on these communities: Tracking (T) becomes substantially more useful when provenance records follow shared standards and can be exported, compared, and re-interpreted by tools other than the one that generated them.




Figure 3: Convergent evolution of empirical scientific practice toward STAMPED. Five independent lines of development are shown on a shared timeline from 1992 to 2025. These comprise foundational ideas, normative frameworks, data version-control tools, platforms and metadata standards, and framework tooling. Each milestone is annotated with the STAMPED principles it anticipates. Despite independent origins, all five tracks converged on the same organizational patterns that STAMPED formalizes. The milestones shown are illustrative rather than exhaustive; many other influential projects and communities contributed to this convergence but are omitted here for brevity.

3.12.3 Adoption incentives and tool certification

STAMPED compliance today is self-assessed against the checklist fulfilling a simple LinkML model²⁰ presented in this paper, and exposed to users in a convenient Web UI at the checklist.stamped-principles.org webpage. That Web UI provides Self-contained (S) stateful permalinks to render any particular state of the filled out checklist, so they could be shared easily. Broader adoption will depend on promotion and external incentives, including funder mandates (e.g., NIH Data Management and Sharing policies), journal reproducibility badges⁴⁷, and venue-level review checklists that converge toward automated scoring. Such a trajectory would parallel that of FAIR, which has moved from principles to machine-assessable indicators with tools such as F-UJI⁴⁸. No automated tool yet available as of now certifies STAMPED compliance. A community-maintained benchmark suite, seeded by the satellite examples project associated with this paper (examples.stamped-principles.org), is a natural precursor.

3.12.4 Teaching and onboarding

Uptake through formal training is slow but compounding. As the “train the trainer” model used by the ReproNim program has demonstrated, students and researchers who encounter these practices early in their career become some of the most effective propagators of open science, even if that propagation takes years to surface. Existing resources already teach most of what STAMPED asks for, without using the term. These include EduWrench⁴⁹, the Reproducible Research MOOC which teaches `git-annex` and related tools⁵⁰, and RDMkit data analysis best practices⁵¹. Curricular integration with open-science and neuroimaging training programs such as Brainhack would close the gap between practitioners who apply these practices by habit and those who can name what they are doing and evaluate it against a shared checklist. YODA principles have already entered the curricula of several such programs, and we expect their references to migrate toward the more formally articulated STAMPED vocabulary. Continued community outreach, including a recent introduction to STAMPED⁵² delivered to the BRAIN BBQS program, supports that transition.

3.12.5 Research in an AI Era

AI agents are now generating analysis code, selecting methods, and in some systems running entire research cycles end-to-end^{53–55}. Without explicit structure, this automation can produce results whose provenance and reasoning are irrecoverable⁵⁶. The STAMPED principles already carry most of what is needed. Self-contained (S) workspaces alone give agents a well-defined boundary; Tracked (T) histories make agent contributions auditable and contribute to their “memory”; Actionable (A) documentation serves humans and machines alike. As with any new methodology or instrumentation entering science, several of these principles need AI-era extensions. Because the agentic-AI ecosystem is evolving rapidly, the directions below are indicative rather than exhaustive.

Self-contained models. Machine learning models are becoming standard components of research workflows, yet they are often fetched at runtime from external repositories with no guarantee of long-term availability or version stability. Users frequently interact with models through platforms without knowing which model version produced their results. Even where models are nominally versioned, the degree of control varies widely, with Hugging Face providing commit-level pinning while commercial APIs may silently update the model behind a fixed endpoint. Self-containment (S) extends naturally to this case: models are Tracked (T) components of the research object, just as data and code are today. Practical mechanisms exist for open models, such as pinned revisions on Hugging Face or, for durable Self-containment (S), archival copies managed via `git-annex`. Tracking what training data and procedures produced a given model remains an open problem that intersects copyright, benchmarking, and provenance, and is mostly beyond what STAMPED can prescribe. Community artifacts such as Model Cards⁵⁷ and Datasheets for Datasets⁵⁸ nonetheless point toward the shape of an answer.

Provenance for AI-driven changes. When an AI agent modifies code, generates analysis scripts, or makes architectural decisions, those actions become part of the research record. Tracking (T) naturally extends to cover agent actions by wrapping an agent invocation in a provenance-capturing command such as `datalad run` to record the model, prompt, and resulting changes alongside the output. Some of this

provenance is already captured organically through git commits authored by agents, session transcripts, and co-authorship metadata. Emerging tools such as Entire CLI⁵⁹ go further, automating capture of the entire AI session alongside the commit. A practical tension follows wherein agent sessions can produce enormous transcripts, and including them wholesale may dilute the value of more targeted provenance records, so what to capture and at what granularity remain open questions. Bitwise reproducibility is not the target and rarely is the goal even for human-driven steps. Instead, what the record must preserve is the hypothesis, the method specification, and enough of the decision trail to support introspection. Self-contained (S) models (above) strengthen this chain—when the model itself is Tracked (T), the provenance record is more complete.

Machine-actionable instructions. Actionability (A) was defined with human operators in mind, including entry-point commands documented in READMEs, frequently incomplete or out of sync with the underlying workflow. AI agents are a second consumer of the same prose, one that can parse, follow, and correct instructions on the fly, and plain Markdown has effectively become the interface language: AGENTS.md files, README-driven agent workflows, and spec-kit⁶⁰ collections now function as machine-actionable interfaces, making Self-contained (S) and Tracked (T) research objects more Actionable (A). Where prose instructions fail, the failure is legible (an audit mechanism that did not exist before) even if agent-driven execution remains less reliable than deterministic tooling.

AI-accelerated specification-to-implementation. Specification-centric research objects (above) depend on the cost of regenerating an implementation from a specification being low enough to be routine. AI coding assistants and specification-driven toolkits such as spec-kit move this cost toward zero, making the spec-centric view practical rather than aspirational. This holds provided the preceding extensions for models, provenance, and machine-actionable instructions are in place. This closes a loop with the end-to-end “AI scientist” systems that opened this section^{53–55}. As soon as an agentic pipeline takes a study specification and returns data, code, figures, and a draft manuscript, the specification and every downstream artifact must together satisfy STAMPED, not only to keep such pipelines efficient to re-run at lower cost but to make their outputs introspectable. Tracked (T) and Self-contained (S) research objects let a researcher inspect which model, prompt, and decision led to a particular result, implement “hand over” between humans and agents alike and let subsequent actor re-enter any prior state and continue from it without rediscovering context. At the end, human scientists authoring and iterating over the specification of a STAMPED research object would see how their edits propagate through the generated implementation. Without the STAMPED principles, the acceleration might be real but the accountability would be lost. Instead, spec-driven agentic research becomes a collaborative, transparent artifact between scientists and the AI systems that support them.

4 Data Availability

All materials, data and code alike, are shared as part of the github.com/stamped-principles GitHub organization, and released under OSI-approved licenses, with per-file licensing metadata maintained via the REUSE specification.

The STAMPED principles⁶¹ (github.com/stamped-principles/stamped-principles-schema) and their accompanying compliance checklist²⁰ (github.com/stamped-principles/stamped-checklist-schema) are formally encoded as LinkML schemas with corresponding JSON instances, separating the data model from the principle content so that each can evolve independently. All sources for this manuscript (e.g., LaTeX file, figures; github.com/stamped-principles/stamped-paper) have a rendered version at paper.stamped-principles.org. Examples illustrating the principles are published as a Hugo-built static website and collected in github.com/stamped-principles/stamped-examples, open for submissions from the community. Examples in that repository are metadata-tagged to allow classification and filtering by principle, domain, and tooling as presented at examples.stamped-principles.org. Visual identity materials (e.g., logos, favicons, etc) are maintained at github.com/stamped-principles/stamped-branding, and organization-level landing content is served from github.com/stamped-principles/ at stamped-principles.org. The bibliography for this work is curated in the Center for Open Neuroscience public Zotero group library at www.zotero.org/groups/6197458/centerforopenneuroscience/library, which is openly reusable.

5 Code Availability

This work is a documentation and formalization project, and the production of scientific software is not its primary aim. The auxiliary code that supports it is distributed across the organization’s repositories and includes tooling around the LinkML schemas in `stamped-principles-schema` and `stamped-checklist-schema`; the front-end JavaScript that powers the interactive compliance checklist deployed at checklist.stamped-principles.org; the Hugo-based site that renders worked examples in `stamped-examples`; and continuous-integration workflows in `stamped-paper` that build the manuscript PDF and post diff previews on pull requests.

References

- [1] Mark D. Wilkinson, Michel Dumontier, IJsbrand Jan Aalbersberg, Gabrielle Appleton, Myles Axton, Arie Baak, Niklas Blomberg, Jan-Willem Boiten, Luiz Bonino Da Silva Santos, Philip E. Bourne, Jildau Bouwman, Anthony J. Brookes, Tim Clark, Mercè Crosas, Ingrid Dillo, Olivier Dumon, Scott Edmunds, Chris T. Evelo, Richard Finkers, Alejandra Gonzalez-Beltran, Alasdair J.G. Gray, Paul Groth, Carole Goble, Jeffrey S. Grethe, Jaap Heringa, Peter A.C ’T Hoen, Rob Hooft, Tobias Kuhn, Ruben Kok, Joost Kok, Scott J. Lusher, Maryann E. Martone, Albert Mons, Abel L. Packer, Bengt Persson, Philippe Rocca-Serra, Marco Roos, Rene Van Schaik, Susanna-Assunta Sansone, Erik Schultes, Thierry Sengstag, Ted Slater, George Strawn, Morris A. Swertz, Mark Thompson, Johan Van Der Lei, Erik Van Mulligen, Jan Velterop, Andra Waagmeester, Peter Wittenburg, Katherine Wolstencroft, Jun Zhao, and Barend Mons. The FAIR Guiding Principles for scientific data management and stewardship. *Scientific Data*, 3(1):160018, March 2016. ISSN 2052-4463. doi: 10.1038/sdata.2016.18. URL <https://www.nature.com/articles/sdata201618>.
- [2] Michelle Barker, Neil P. Chue Hong, Daniel S. Katz, Anna-Lena Lamprecht, Carlos Martinez-Ortiz, Fotis Psomopoulos, Jennifer Harrow, Leyla Jael Castro, Morane Gruenpeter, Paula Andrea Martinez, and Tom Honeyman. Introducing the FAIR Principles for research software. *Scientific Data*, 9(1):622, October 2022. ISSN 2052-4463. doi: 10.1038/s41597-022-01710-x. URL <https://www.nature.com/articles/s41597-022-01710-x>.
- [3] Sean R. Wilkinson, Meznah Aloqalaa, Khalid Belhajjame, Michael R. Crusoe, Bruno De Paula Kinoshita, Luiz Gadelha, Daniel Garijo, Ove Johan Ragnar Gustafsson, Nick Juty, Sehrish Kanwal, Farah Zaib Khan, Johannes Köster, Karsten Peters-von Gehlen, Line Pouchard, Randy K. Rannow, Stian Soiland-Reyes, Nicola Soranzo, Shoaib Sufi, Ziheng Sun, Baiba Vilne, Merridee A. Wouters, Denis Yuen, and Carole Goble. Applying the FAIR Principles to computational workflows. *Scientific Data*, 12(1):328, February 2025. ISSN 2052-4463. doi: 10.1038/s41597-025-04451-9. URL <https://www.nature.com/articles/s41597-025-04451-9>.
- [4] Michael Hanke, Matteo Visconti di Oleggio Castello, Kyle Meyer, Benjamin Poldrack, and Yaroslav O. Halchenko. YODA: YODA’s organigram on data analysis, 2018. URL <https://github.com/myyoda/poster/blob/master/ohbm2018.pdf>. Published: Poster presented at the annual meeting of the Organization for Human Brain Mapping, Singapore.
- [5] Michael Hanke. What is DataLad and what can it do for you?, 2023. URL https://files.inm7.de/mih/pres/talks/whatisdatalad_2023.html.
- [6] S. Bradner. Key words for use in RFCs to Indicate Requirement Levels. Technical Report RFC2119, RFC Editor, March 1997. URL <https://www.rfc-editor.org/info/rfc2119>.
- [7] Yaroslav O. Halchenko, Kyle Meyer, Benjamin Poldrack, Debanjum Singh Solanky, Adina S. Wagner, Jason Gors, Dave MacFarlane, Dorian Pustina, Vanessa Sochat, Satrajit S. Ghosh, Christian Mönch, Christopher J. Markiewicz, Laura Waite, Ilya Shlyakhter, Alejandro de la Vega, Soichi Hayashi, Christian Olaf Häusler, Jean-Baptiste Poline, Tobias Kadelka, Kusti Skytén, Dorota Jarecka, David Kennedy, Ted Strauss, Matt Cieslak, Peter Vavra, Horea-Ioan Ioanas, Robin Schneider, Mika Pflüger, James V. Haxby, Simon B. Eickhoff, and Michael Hanke. DataLad: distributed system for joint management

- of code, data, and their relationship. *Journal of Open Source Software*, 6(63):3262, July 2021. ISSN 2475-9066. doi: 10.21105/joss.03262. URL <https://joss.theoj.org/papers/10.21105/joss.03262>.
- [8] Ruslan Kuprieiev, skshetry, Peter Rowlands, Dmitry Petrov, Paweł Redzyński, Casper da Costa-Luis, David de la Iglesia Castro, Alexander Schepanovski, Ivan Shcheklein, Yanxiang Gao, Batuhan Taskaya, Jorge Orpinel, Fábio Santos, Daniele, Ronan Lamy, Aman Sharma, Zhanibek Kaimuldenov, Dani Hodovic, Nikita Kodenko, Earl, Andrew Grigorev, Nabanita Dash, Dave Berenbaum, George Vyshnya, maykulkarni, Max Hora, Vera, and Sanidhya Mangal. DVC: Data Version Control - Git for Data & Models, March 2026. URL <https://zenodo.org/doi/10.5281/zenodo.3677553>.
- [9] Erik C. Johnson, Thinh T. Nguyen, Benjamin K. Dichter, Frank Zappulla, Montgomery Kosma, Kabilar Gunalan, Yaroslav O. Halchenko, Shay Q. Neufeld, Kristen Ratan, Nicholas J. Edwards, Susanne Ressler, Sarah R. Heilbronner, Michael Schirner, Petra Ritter, Brock Wester, Satrajit Ghosh, Maryann E. Martone, Franco Pestilli, and Dimitri Yatsenko. SciOps: Achieving Productivity and Reliability in Data-Intensive Research, 2024. URL <https://arxiv.org/abs/2401.00077>.
- [10] W3C Provenance Working Group, Yolanda Gil, and Simon Miles. PROV-Overview: An Overview of the PROV Family of Documents, April 2013. URL <https://www.w3.org/TR/prov-overview/>. Published: W3C Working Group Note.
- [11] Stian Soiland-Reyes, Peter Sefton, Mercè Crosas, Leyla Jael Castro, Frederik Coppens, José M. Fernández, Daniel Garijo, Björn Grüning, Marco La Rosa, Simone Leo, Eoghan Ó Carragáin, Marc Portier, Ana Trisovic, RO-Crate Community, Paul Groth, and Carole Goble. Packaging research artefacts with RO-Crate. *Data Science*, 5(2):97–138, July 2022. ISSN 2451-8484, 2451-8492. doi: 10.3233/DS-210053. URL <https://journals.sagepub.com/doi/10.3233/DS-210053>.
- [12] BIDS Community. BEP028: BIDS Extension Proposal for BIDS Provenance, 2025. URL https://github.com/bids-standard/BEP028_BIDSprov. Published: BIDS Extension Proposal.
- [13] Tek Raj Chhetri, Yaroslav O. Halchenko, Dorota Jarocka, Priya Trivedi, Satrajit Ghosh, Pat Ray, and Lydia Ng. Big Data in Scientific Knowledge Graphs: Reproducibility, A Survey of Provenance, Association, and Evidence Ontologies. In *Proceedings of the 2025 ACM Web Conference 2025 (WWW '25)*, Sydney, NSW, Australia, May 2025. ACM, 2025. ISBN 978-1-60558-331-6. doi: 10.1145/3701716.3715483. URL <https://doi.org/10.1145/3701716.3715483>.
- [14] Markus Kowalewski and Phillip Seeber. Sustainable packaging of quantum chemistry software with the Nix package manager. *International Journal of Quantum Chemistry*, 122(9):e26872, May 2022. ISSN 0020-7608, 1097-461X. doi: 10.1002/qua.26872. URL <https://onlinelibrary.wiley.com/doi/10.1002/qua.26872>.
- [15] Ludovic Courtès and Ricardo Wurmus. Reproducible and User-Controlled Software Environments in HPC with Guix, 2015. URL <https://arxiv.org/abs/1506.02822>. Version Number: 2.
- [16] Adina S. Wagner, Laura K. Waite, Malgorzata Wierzbza, Felix Hoffstaedter, Alexander Q. Waite, Benjamin Poldrack, Simon B. Eickhoff, and Michael Hanke. FAIRly big: A framework for computationally reproducible processing of large-scale data. *Scientific Data*, 9(1):80, March 2022. ISSN 2052-4463. doi: 10.1038/s41597-022-01163-2. URL <https://www.nature.com/articles/s41597-022-01163-2>.
- [17] Barbara Cheifet. Promoting reproducibility with Code Ocean. *Genome Biology*, 22(1):65, s13059–021–02299–x, December 2021. ISSN 1474-760X. doi: 10.1186/s13059-021-02299-x. URL <https://genomebiology.biomedcentral.com/articles/10.1186/s13059-021-02299-x>.
- [18] Kate Stewart, Phil Odenice, and Esteban Rockett. Software package data exchange (SPDX) specification. *International Free and Open Source Software Law Review*, 2(2):191–196, 2010. URL <https://www.ifosslr.org/ifosslr/article/view/45>.
- [19] Free Software Foundation Europe. REUSE specification, version 3.3, 2024. URL <https://reuse.software/spec-3.3/>.

- [20] Cody Baker. stamped-principles/stamped-checklist-schema: v0.1.0, April 2026. URL <https://zenodo.org/doi/10.5281/zenodo.19672391>.
- [21] Yaroslav Halchenko, Austin Macdonald, and Jiameng Wu. stamped-principles/stamped-examples: Pre-print ready, May 2026. URL <https://zenodo.org/doi/10.5281/zenodo.20315397>.
- [22] Krzysztof J. Gorgolewski, Tibor Auer, Vince D. Calhoun, R. Cameron Craddock, Samir Das, Eugene P. Duff, Guillaume Flandin, Satrajit S. Ghosh, Tristan Glatard, Yaroslav O. Halchenko, Daniel A. Handwerker, Michael Hanke, David Keator, Xiangrui Li, Zachary Michael, Camille Maumet, B. Nolan Nichols, Thomas E. Nichols, John Pellman, Jean-Baptiste Poline, Ariel Rokem, Gunnar Schaefer, Vanessa Sochat, William Triplett, Jessica A. Turner, Gaël Varoquaux, and Russell A. Poldrack. The brain imaging data structure, a format for organizing and describing outputs of neuroimaging experiments. *Scientific Data*, 3(1):000044, June 2016. ISSN 2052-4463. doi: 10.1038/sdata.2016.44. URL <https://www.nature.com/articles/sdata201644>.
- [23] Oscar Esteban, Rastko Ciric, Karolina Finc, Ross Blair, Christopher J. Markiewicz, Craig A. Moodie, James D. Kent, Mathias Goncalves, Elizabeth DuPre, Daniel E. P. Gomez, Zhifang Ye, Taylor Salo, Romain Valabregue, Inge K. Amlien, Franziskus Liem, Nir Jacoby, Hrvoje Stojić, Matthew Cieslak, Sebastian Urchs, Yaroslav O. Halchenko, Satrajit S. Ghosh, Alejandro De La Vega, Tal Yarkoni, Jesse Wright, William H. Thompson, Russell A. Poldrack, and Krzysztof J. Gorgolewski. Analysis of task-based functional MRI data preprocessed with fMRIPrep. *bioRxiv*, 2020. doi: 10.1101/694364. URL <https://www.biorxiv.org/content/early/2020/01/03/694364>.
- [24] Christopher J Markiewicz, Krzysztof J Gorgolewski, Franklin Feingold, Ross Blair, Yaroslav O Halchenko, Eric Miller, Nell Hardcastle, Joe Wexler, Oscar Esteban, Mathias Goncalves, Anita Jwa, and Russell Poldrack. The OpenNeuro resource for sharing of neuroscience data. *eLife*, 10:e71774, October 2021. ISSN 2050-084X. doi: 10.7554/eLife.71774. URL <https://doi.org/10.7554/eLife.71774>.
- [25] Chenying Zhao, Dorota Jarecka, Sydney Covitz, Yibei Chen, Simon B. Eickhoff, Damien A. Fair, Alexandre R. Franco, Yaroslav O. Halchenko, Timothy J. Hendrickson, Felix Hoffstaedter, Audrey Houghton, Gregory Kiar, Austin Macdonald, Kahini Mehta, Michael P. Milham, Taylor Salo, Michael Hanke, Satrajit S. Ghosh, Matthew Cieslak, and Theodore D. Satterthwaite. A reproducible and generalizable software workflow for analysis of large-scale neuroimaging data collections using BIDS Apps. *Imaging Neuroscience*, 2:imag-2-00074, January 2024. ISSN 2837-6056. doi: 10.1162/imag-a-00074. URL <https://direct.mit.edu/imag/article/doi/10.1162/imag-a-00074/119046/A-reproducible-and-generalizable-software-workflow>.
- [26] Jon F. Claerbout and Martin Karrenbach. Electronic documents give reproducible research a new meaning. In *SEG Technical Program Expanded Abstracts 1992*, pages 601–604. Society of Exploration Geophysicists, January 1992. doi: 10.1190/1.1822162. URL <http://library.seg.org/doi/abs/10.1190/1.1822162>.
- [27] Jonathan B. Buckheit and David L. Donoho. WaveLab and Reproducible Research. In P. Bickel, P. Diggle, S. Fienberg, K. Krickeberg, I. Olkin, N. Wermuth, S. Zeger, Anestis Antoniadis, and Georges Oppenheim, editors, *Wavelets and Statistics*, volume 103, pages 55–81. Springer New York, New York, NY, 1995. ISBN 978-0-387-94564-4 978-1-4612-2544-7. doi: 10.1007/978-1-4612-2544-7_5. URL http://link.springer.com/10.1007/978-1-4612-2544-7_5. Series Title: Lecture Notes in Statistics.
- [28] Robert Gentleman and Duncan Temple Lang. Statistical Analyses and Reproducible Research. *Journal of Computational and Graphical Statistics*, 16(1):1–23, March 2007. ISSN 1061-8600, 1537-2715. doi: 10.1198/106186007X178663. URL <http://www.tandfonline.com/doi/abs/10.1198/106186007X178663>.
- [29] William Stafford Noble. A Quick Guide to Organizing Computational Biology Projects. *PLoS Computational Biology*, 5(7):e1000424, July 2009. ISSN 1553-7358. doi: 10.1371/journal.pcbi.1000424. URL <https://dx.plos.org/10.1371/journal.pcbi.1000424>.
- [30] Victoria Stodden. The Scientific Method in Practice: Reproducibility in the Computational Sciences. *SSRN Electronic Journal*, 2010. ISSN 1556-5068. doi: 10.2139/ssrn.1550193. URL <http://www.ssrn.com/abstract=1550193>.

- [31] Geir Kjetil Sandve, Anton Nekrutenko, James Taylor, and Eivind Hovig. Ten Simple Rules for Reproducible Computational Research. *PLoS Computational Biology*, 9(10):e1003285, October 2013. ISSN 1553-7358. doi: 10.1371/journal.pcbi.1003285. URL <https://dx.plos.org/10.1371/journal.pcbi.1003285>.
- [32] Greg Wilson, Jennifer Bryan, Karen Cranston, Justin Kitzes, Lex Nederbragt, and Tracy K. Teal. Good enough practices in scientific computing. *PLOS Computational Biology*, 13(6):e1005510, June 2017. ISSN 1553-7358. doi: 10.1371/journal.pcbi.1005510. URL <https://dx.plos.org/10.1371/journal.pcbi.1005510>.
- [33] Koenraad De Smedt, Dimitris Koureas, and Peter Wittenburg. FAIR Digital Objects for Science: From Data Pieces to Actionable Knowledge Units. *Publications*, 8(2):21, April 2020. ISSN 2304-6775. doi: 10.3390/publications8020021. URL <https://www.mdpi.com/2304-6775/8/2/21>.
- [34] Ryan Abernathy. Icechunk: An Open-Source Transactional Storage Engine for Zarr. October 2024. doi: 10.5281/ZENODO.13989108. URL <https://zenodo.org/doi/10.5281/zenodo.13989108>.
- [35] Enis Afgan, Dannon Baker, Bérénice Batut, Marius van den Beek, Dave Bouvier, Martin Čech, John Chilton, Dave Clements, Nate Coraor, Björn A Gruning, Aysam Guerler, Jennifer Hillman-Jackson, Saskia Hiltemann, Vahid Jalili, Helena Rasche, Nicola Soranzo, Jeremy Goecks, James Taylor, Anton Nekrutenko, and Daniel Blankenberg. The Galaxy platform for accessible, reproducible and collaborative biomedical analyses: 2018 update. *Nucleic Acids Research*, 46(W1):W537–W544, July 2018. ISSN 0305-1048, 1362-4962. doi: 10.1093/nar/gky379. URL <https://academic.oup.com/nar/article/46/W1/W537/5001157>.
- [36] Soichi Hayashi, Bradley A. Caron, Anibal Sólón Heinsfeld, Sophia Vinci-Booher, Brent McPherson, Daniel N. Bullock, Giulia Bertò, Guiomar Niso, Sandra Hanekamp, Daniel Levitas, Kimberly Ray, Anne MacKenzie, Paolo Avesani, Lindsey Kitchell, Josiah K. Leong, Filipi Nascimento-Silva, Serge Koudoro, Hanna Willis, Jasleen K. Jolly, Derek Pisner, Taylor R. Zuidema, Jan W. Kurzawski, Kyriaki Mikellidou, Aurore Bussalbé, Maximilien Chaumon, Nathalie George, Christopher Rorden, Conner Victory, Dheeraj Bhatia, Dogu Baran Aydogan, Fang-Cheng F. Yeh, Franco Delogu, Javier Guaje, Jelle Veraart, Jeremy Fischer, Joshua Faskowitz, Ricardo Fabrega, David Hunt, Shawn McKee, Shawn T. Brown, Stephanie Heyman, Vittorio Iacovella, Amanda F. Mejia, Daniele Marinazzo, R. Cameron Craddock, Emanuele Olivetti, Jamie L. Hanson, Eleftherios Garyfallidis, Dan Stanzione, James Carson, Robert Henschel, David Y. Hancock, Craig A. Stewart, David Schnyer, Damian O. Eke, Russell A. Poldrack, Steffen Bollmann, Ashley Stewart, Holly Bridge, Ilaria Sani, Winrich A. Freiwald, Aina Puce, Nicholas L. Port, and Franco Pestilli. brainlife.io: a decentralized and open-source cloud platform to support neuroscience research. *Nature Methods*, 21(5):809–813, May 2024. ISSN 1548-7091, 1548-7105. doi: 10.1038/s41592-024-02237-2. URL <https://www.nature.com/articles/s41592-024-02237-2>.
- [37] Dan Fowler, Jo Barratt, and Paul Walsh. Frictionless Data: Making Research Data Quality Visible. *International Journal of Digital Curation*, 12(2):274–285, May 2018. ISSN 1746-8256. doi: 10.2218/ijdc.v12i2.577. URL <http://ijdc.net/article/view/577>.
- [38] Nicole Cendana Ferrer, Janisha Patel, Dennis A. Dean, Jonathon Keeney, Jianchao Yao, Jeremy Goecks, Stian Soiland-Reyes, Paul Duncan, Robel Kahsay, Anais Hayes, Michael R. Crusoe, Charles Hadley King, Carole Goble, Amanda Bell, Gil Alterovitz, Eric Donaldson, Vahan Simonyan, Jonas Almeida, Mark Walderhaug, Elaine Thompson, Krista Smith, and Raja Mazumder. BioCompute Project (#BioComputeObjects). 2017. doi: 10.17605/OSF.IO/H59UH. URL <https://osf.io/h59uh/>.
- [39] Carole Goble, Sarah Cohen-Boulakia, Stian Soiland-Reyes, Daniel Garijo, Yolanda Gil, Michael R. Crusoe, Kristian Peters, and Daniel Schober. FAIR Computational Workflows. *Data Intelligence*, 2(1-2): 108–121, January 2020. ISSN 2641-435X. doi: 10.1162/dint_a_00033. URL https://www.sciengine.com/doi/10.1162/dint_a_00033.
- [40] Brian A. Nosek, Charles R. Ebersole, Alexander C. DeHaven, and David T. Mellor. The preregistration revolution. *Proceedings of the National Academy of Sciences*, 115(11):2600–2606, March 2018. ISSN

- 0027-8424, 1091-6490. doi: 10.1073/pnas.1708274114. URL <https://pnas.org/doi/full/10.1073/pnas.1708274114>.
- [41] Erin D. Foster and Ariel Deardorff. Open Science Framework (OSF). *Journal of the Medical Library Association*, 105(2), April 2017. ISSN 1558-9439, 1536-5050. doi: 10.5195/jmla.2017.88. URL <http://jmla.pitt.edu/ojs/jmla/article/view/88>.
 - [42] Michelle Wang, Nikhil Bhagwat, Francesco Cremonesi, Mathieu Dugré, Julia-Katharina Pfarr, Emile d’Angremont, Alyssa Dai, Arman Jahanpour, Sebastian Urchs, Sergen Cansiz, Lucie Chambon, Ali Tolga Dinger, Jhonatan Torres, Marc Vesin, Gabriel Pinilla-Monsalve, Yuan Song, Chris Vriend, Francis Jeanson, Oury Monchi, Ysbrand D. Van Der Werf, Marco Lorenzi, and Jean-Baptiste Poline. Open neuroinformatics infrastructure ecosystem for federated multisite studies, May 2026. URL <http://biorxiv.org/lookup/doi/10.64898/2026.04.30.721944>.
 - [43] Michael R. Crusoe, Sanne Abeln, Alexandru Iosup, Peter Amstutz, John Chilton, Nebojša Tijanić, Hervé Ménager, Stian Soiland-Reyes, Bogdan Gavrilović, Carole Goble, and The Cwl Community. Methods included: standardizing computational reuse and portability with the Common Workflow Language. *Communications of the ACM*, 65(6):54–63, June 2022. ISSN 0001-0782, 1557-7317. doi: 10.1145/3486897. URL <https://dl.acm.org/doi/10.1145/3486897>.
 - [44] Felix Mölder, Kim Philipp Jablonski, Brice Letcher, Michael B. Hall, Christopher H. Tomkins-Tinch, Vanessa Sochat, Jan Forster, Soohyun Lee, Sven O. Twardziok, Alexander Kanitz, Andreas Wilm, Manuel Holtgrewe, Sven Rahmann, Sven Nahnsen, and Johannes Köster. Sustainable data analysis with Snakemake. *F1000Research*, 10:33, April 2021. ISSN 2046-1402. doi: 10.12688/f1000research.29032.2. URL <https://f1000research.com/articles/10-33/v2>.
 - [45] Paolo Di Tommaso, Maria Chatzou, Evan W Floden, Pablo Prieto Barja, Emilio Palumbo, and Cedric Notredame. Nextflow enables reproducible computational workflows. *Nature Biotechnology*, 35(4): 316–319, April 2017. ISSN 1087-0156, 1546-1696. doi: 10.1038/nbt.3820. URL <https://www.nature.com/articles/nbt.3820>.
 - [46] Rotem Botvinik-Nezer, Felix Holzmeister, Colin F. Camerer, Anna Dreber, Juergen Huber, Magnus Johannesson, Michael Kirchler, Roni Iwanir, Jeanette A. Mumford, R. Alison Adcock, Paolo Avesani, Blazej M. Baczkowski, Aahana Bajracharya, Leah Bakst, Sheryl Ball, Marco Barilari, Nadège Bault, Derek Beaton, Julia Beitner, Roland G. Benoit, Ruud M. W. J. Berkers, Jamil P. Bhanji, Bharat B. Biswal, Sebastian Bobadilla-Suarez, Tiago Bortolini, Katherine L. Bottenhorn, Alexander Bowring, Senne Braem, Hayley R. Brooks, Emily G. Brudner, Cristian B. Calderon, Julia A. Camilleri, Jaime J. Castrellon, Luca Cecchetti, Edna C. Cieslik, Zachary J. Cole, Olivier Collignon, Robert W. Cox, William A. Cunningham, Stefan Czoschke, Kamalaker Dadi, Charles P. Davis, Alberto De Luca, Mauricio R. Delgado, Lysia Demetriou, Jeffrey B. Dennison, Xin Di, Erin W. Dickie, Ekaterina Dobryakova, Claire L. Donnat, Juergen Dukart, Niall W. Duncan, Joke Durnez, Amr Eed, Simon B. Eickhoff, Andrew Erhart, Laura Fontanesi, G. Matthew Fricke, Shiguang Fu, Adriana Galván, Remi Gau, Sarah Genon, Tristan Glatard, Enrico Glerean, Jelle J. Goeman, Sergej A. E. Golwin, Carlos González-García, Krzysztof J. Gorgolewski, Cheryl L. Grady, Mikella A. Green, João F. Guassi Moreira, Olivia Guest, Shabnam Hakimi, J. Paul Hamilton, Roeland Hancock, Giacomo Handjaras, Bronson B. Harry, Colin Hawco, Peer Herholz, Gabrielle Herman, Stephan Heunis, Felix Hoffstaedter, Jeremy Hogeveen, Susan Holmes, Chuan-Peng Hu, Scott A. Huettel, Matthew E. Hughes, Vittorio Iacovella, Alexandru D. Iordan, Peder M. Isager, Ayse I. Isik, Andrew Jahn, Matthew R. Johnson, Tom Johnstone, Michael J. E. Joseph, Anthony C. Juliano, Joseph W. Kable, Michalis Kassinosopoulos, Cemal Koba, Xiang-Zhen Kong, Timothy R. Koscik, Nuri Erkut Kucukboyaci, Brice A. Kuhl, Sebastian Kupek, Angela R. Laird, Claus Lamm, Robert Langner, Nina Lauharatanahirun, Hongmi Lee, Sangil Lee, Alexander Leemans, Andrea Leo, Elise Lesage, Flora Li, Monica Y. C. Li, Phui Cheng Lim, Evan N. Lintz, Schuyler W. Liphardt, Annabel B. Losecaat Vermeer, Bradley C. Love, Michael L. Mack, Norberto Malpica, Theo Marins, Camille Maumet, Kelsey McDonald, Joseph T. McGuire, Helena Melero, Adriana S. Méndez Leal, Benjamin Meyer, Kristin N. Meyer, Glad Mihai, Georgios D. Mitsis, Jorge Moll, Dylan M. Nielson, Gustav Nilsson, Michael P. Notter, Emanuele Olivetti, Adrian I. Onicas, Paolo Papale, Kaustubh R. Patil, Jonathan E.

- Peelle, Alexandre Pérez, Doris Pischedda, Jean-Baptiste Poline, Yanina Prystauka, Shruti Ray, Patricia A. Reuter-Lorenz, Richard C. Reynolds, Emiliano Ricciardi, Jenny R. Rieck, Anaïs M. Rodriguez-Thompson, Anthony Romy, Taylor Salo, Gregory R. Samanez-Larkin, Emilio Sanz-Morales, Margaret L. Schlichting, Douglas H. Schultz, Qiang Shen, Margaret A. Sheridan, Jennifer A. Silvers, Kenny Skagerlund, Alec Smith, David V. Smith, Peter Sokol-Hessner, Simon R. Steinkamp, Sarah M. Tashjian, Bertrand Thirion, John N. Thorp, Gustav Tinghög, Loreen Tisdall, Steven H. Thompson, Claudio Toro-Serey, Juan Jesus Torre Tresols, Leonardo Tozzi, Vuong Truong, Luca Turella, Anna E. Van 'T Veer, Tom Verguts, Jean M. Vettel, Sagana Vijayarajah, Khoi Vo, Matthew B. Wall, Wouter D. Weeda, Susanne Weis, David J. White, David Wisniewski, Alba Xifra-Porxas, Emily A. Yearling, Sangsuk Yoon, Rui Yuan, Kenneth S. L. Yuen, Lei Zhang, Xu Zhang, Joshua E. Zosky, Thomas E. Nichols, Russell A. Poldrack, and Tom Schonberg. Variability in the analysis of a single neuroimaging dataset by many teams. *Nature*, 582(7810):84–88, June 2020. ISSN 0028-0836, 1476-4687. doi: 10.1038/s41586-020-2314-9. URL <https://www.nature.com/articles/s41586-020-2314-9>.
- [47] Center for Open Science. Open Science Badges, 2023. URL <https://www.cos.io/initiatives/badges>. Published: Center for Open Science initiative.
- [48] Anusuriya Devaraju and Robert Huber. An automated solution for measuring the progress toward FAIR research data. *Patterns*, 2(11):100370, November 2021. ISSN 2666-3899. doi: 10.1016/j.patter.2021.100370. URL <https://doi.org/10.1016/j.patter.2021.100370>.
- [49] Henri Casanova and others. EduWrench: Pedagogic Modules for Parallel and Distributed Computing, 2023. URL <https://eduwrench.ics.hawaii.edu/>. Published: Online teaching platform, University of Hawai‘i.
- [50] Arnaud Legrand, Martin Quinson, Bruno Bzeznik, Vincent Danjean, Michael Mercier, and Olivier Richard. Reproducible Research II: Practices and Tools for Managing Computations and Data, 2022. URL <https://www.fun-mooc.fr/en/courses/reproducible-research-ii-practices-and-tools-for-managing-comput/>. Published: FUN-MOOC online course.
- [51] ELIXIR RDMkit Community. Data analysis best practices, 2024. URL https://rdmkit.elixir-europe.org/data_analysis. Published: ELIXIR Research Data Management Kit.
- [52] Cody Baker. Guidelines for Reproducible Research (STAMPED), March 2026. URL <https://www.youtube.com/watch?v=8NTWKHer5Zo>. Published: Talk presented at the Virtual BBQS Workshop, Center for Open Neuroscience.
- [53] Daniil A. Boiko, Robert MacKnight, Ben Kline, and Gabe Gomes. Autonomous chemical research with large language models. *Nature*, 624(7992):570–578, December 2023. ISSN 0028-0836, 1476-4687. doi: 10.1038/s41586-023-06792-0. URL <https://www.nature.com/articles/s41586-023-06792-0>.
- [54] Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery, September 2024. URL <http://arxiv.org/abs/2408.06292>. arXiv:2408.06292 [cs].
- [55] Ludovico Mitchener, Angela Yiu, Benjamin Chang, Mathieu Bourdenx, Tyler Nadolski, Arvis Sulovari, Eric C. Landsness, Daniel L. Barabasi, Siddharth Narayanan, Nicky Evans, Shriya Reddy, Martha Foiani, Aizad Kamal, Leah P. Shriver, Fang Cao, Asmamaw T. Wassie, Jon M. Laurent, Edwin Melville-Green, Mayk Caldas, Albert Bou, Kaleigh F. Roberts, Sladjana Zagorac, Timothy C. Orr, Miranda E. Orr, Kevin J. Zvezdaryk, Ali E. Ghareeb, Laurie McCoy, Bruna Gomes, Euan A. Ashley, Karen E. Duff, Tonio Buonassisi, Tom Rainforth, Randall J. Bateman, Michael Skarlinski, Samuel G. Rodrigues, Michaela M. Hinks, and Andrew D. White. Kosmos: An AI Scientist for Autonomous Discovery, November 2025. URL <http://arxiv.org/abs/2511.02824>. arXiv:2511.02824 [cs].
- [56] Lisa Messeri and M. J. Crockett. Artificial intelligence and illusions of understanding in scientific research. *Nature*, 627(8002):49–58, March 2024. ISSN 0028-0836, 1476-4687. doi: 10.1038/s41586-024-07146-0. URL <https://www.nature.com/articles/s41586-024-07146-0>.

- [57] Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. Model Cards for Model Reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency*, pages 220–229, Atlanta GA USA, January 2019. ACM. ISBN 978-1-4503-6125-5. doi: 10.1145/3287560.3287596. URL <https://dl.acm.org/doi/10.1145/3287560.3287596>.
- [58] Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé Iii, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12):86–92, December 2021. ISSN 0001-0782, 1557-7317. doi: 10.1145/3458723. URL <https://dl.acm.org/doi/10.1145/3458723>.
- [59] Entire. Entire CLI: Git observability layer that captures AI agent sessions as checkpoints alongside commits, 2026. URL <https://github.com/entireio/cli>. Published: Software, GitHub repository.
- [60] GitHub. Spec Kit: Toolkit for Spec-Driven Development, 2025. URL <https://github.com/github/spec-kit>. Published: Software, GitHub repository.
- [61] Cody Baker. stamped-principles/stamped-principles-schema: v0.1.0, April 2026. URL <https://zenodo.org/doi/10.5281/zenodo.19672387>.

6 Author Contributions

A.M.: Conceptualization, Methodology, Software, Resources, Writing – original draft, and Writing – review & editing. C.C.B.: Conceptualization, Methodology, Software, Visualization, Writing – original draft, and Writing – review & editing. I.T.: Writing – review & editing. Y.O.H.: Conceptualization, Funding acquisition, Supervision (lead), Project administration, Methodology, Software, Resources, and Writing – review & editing.

7 Competing Interests

The authors declare no competing interests.

8 Acknowledgments

STAMPED builds directly on prior community efforts to articulate principles for reproducible research data management. We are grateful to Michael Hanke and Adina S. Wagner for the conceptual development of VAMP⁵, and to the authors of the original YODA poster⁴—Michael Hanke, Matteo Visconti di Oleggio Castello, Kyle Meyer, and Benjamin Poldrack—whose work laid the foundations on which this formalization rests. We thank Jiameng Wu for filing issues and contributing editorial improvements to the manuscript.

This manuscript and companion materials in the github.com/stamped-principles organization were prepared with the assistance of agentic AI coding tools and large language models (primarily Claude Opus). All final text, figures, and specifications were edited and/or confirmed by the human authors, who are responsible for the content. Where AI tools notably contributed to a change, we have strived to annotate the corresponding commits with a Co-Authored-By trailer, or analogous comments, to identify the tool and potentially the model version, so that the provenance of individual contributions is inspectable in the public git history of the `stamped-*` repositories.

9 Funding

This work was supported by the National Institutes of Health: ReproNim — Center for Reproducible Neuroimaging Computation (NIBIB P41 EB019936), OpenNeuro — An Open Archive for Analysis and Sharing of BRAIN Initiative Data (NIMH R24 MH117179), DANDI — Distributed Archives for Neurophysiology

Data Integration (NIMH R24 MH117295), and EMBER — Ecosystem for Multi-modal Brain-behavior Experimentation and Research (NIMH R24 MH136632).